

SHASHANK PRASANNA

shashankprasanna.com

(480) 347 5005 • shashank4@gmail.com

PROFILE

- Developer Relations leader with 15+ years building AI/ML developer ecosystems at Apple, NVIDIA, AWS, Meta, and Modular. Owned launch strategy, content, and adoption across NVIDIA's AI/ML software stack (cuDNN, TensorRT, RAPIDS, GPU containers), PyTorch 2.0, Mojo and MAX, and Apple's MLX, Core ML, Core AI, and Foundation Models frameworks.
- Build what I teach: private and secure AI use cases including agentic RAG applications on Apple platforms with multi-phase tool-calling loops, hybrid retrieval, sandboxed code execution, and swappable on-device, local, and Claude backends; evaluated against failure-mode datasets with Apple's Evaluations framework. Drive agentic development with Claude Code and Claude Agent in Xcode.
- Scale through people and community – coached 15+ engineering presenters, hired and led content contributors, ran NVIDIA's deep learning framework advisory board, grew reach through AWS Heroes. 100+ talks and articles; 500K+ views on recent work. Portfolio: **shashankprasanna.com**.

WORK EXPERIENCE

Apple - AI/ML Technologies Evangelism

Aug 2024 - present

Lead AI/ML developer messaging, content strategy, and adoption programs for Apple registered developers. Own WWDC AI/ML content planning across Foundation Models, MLX, and on-device ML APIs; coached 15 engineering presenters, including content guidance for a WWDC26 session at 440K+ views. Delivered talks, panels, and on-camera code-alongs at 265K+ views, including a public code-along teaching Apple's on-device LLM API. Content links on shashankprasanna.com.

Provide guidance to 20–30 top-priority App Store and indie developers on what features to build with on-device LLM and ML frameworks and how to build them. Consult with developers to take shipping AI features from good to great. Direct architecture across agentic assistants, agentic retrieval-augmented pipelines, and on-device inference performance; advise on choosing among model providers, from Apple's on-device model to server-side models like Claude, and on agentic development in Xcode with coding agents like Claude Agent. Carry field feedback into framework and API planning.

Modular - AI Developer Advocacy & Developer Marketing Tech Lead

June 2023 - Aug 2024

Owned developer marketing and advocacy in an engineering-heavy AI startup, building awareness and community for Mojo and MAX before broad commercialization. Defined and executed launch marketing, developer content, and community growth plans. Authored and edited Modular developer blogs; managed Modular YouTube content; created and hosted Modular Community Q&A Livestreams; and ran hands-on workshops. Served as team lead, helped hire marketing contributors, and built low-level programming fluency to create credible technical content.

Meta - PyTorch and ML Open-Source Developer Advocate

Sept 2022 - June 2023

Led ML OSS and PyTorch developer engagement for the PyTorch 2.0 launch. Developed and executed launch content strategy; planned and hosted PyTorch 2.0 Ask the Engineers livestreams with PyTorch and partner engineers; partnered with the Linux Foundation and PyTorch Foundation members (AWS, Microsoft) on developer training; and authored a widely read PyTorch 2.0 technical deep dive.

Amazon Web Services - AI/ML Technical Evangelist

Jan 2019 - Sept 2022

Led North America AI/ML developer outreach and education across open-source ML frameworks, containers, GPUs, Intel Habana Gaudi, AWS Inferentia & Trainium, Amazon SageMaker, and EKS. Built awareness and consideration through 10+ workshops and bootcamps, often with 300-500 live attendees, plus partner programs

with Intel, Meta, and NVIDIA. Scaled reach through AWS Heroes and builder communities and influenced product roadmap as the voice of developers.

NVIDIA - *Technical Developer Product Marketing, Deep Learning* *Jan 2016 - Jan 2019*

First product marketing manager focused on deep learning software at NVIDIA. Established developer marketing programs for GPU-accelerated deep learning and helped grow the community from early adoption. Launched and positioned cuDNN, NCCL, TensorRT, and GPU support for Containers/Kubernetes; created benchmarks and technical webpages; and shaped library strategy by running the deep learning framework advisory board (PyTorch, TensorFlow, MXNet, Chainer, Caffe, Theano, more). Editor and author for the NVIDIA developer blog.

MATHWORKS - *Product Manager, Statistics & Machine Learning* *Sep 2011 - Jan 2016*

Sep 2013 - Jan 2016: Owned product strategy for statistics, machine learning and data analytics business areas. Led the development of product roadmaps, annual marketing plans, and launch plans. Created developer-focused content such as webinars and demos (1) Statistics and Machine Learning Toolbox (2) Deep Learning Toolbox (3) Curve Fitting Toolbox.

Jul 2011 - Aug 2013: Application engineer responsible for sales support projects for the MathWorks Machine Learning, Statistical Analysis, Data Analytics and Computational Finance product areas.

ORACLE - *Software Developer* *May 2008 - July 2009*

C++ software developer responsible for product design and development projects for Oracle Siebel CRM

EDUCATION

Arizona State University *July 2011*

M.S. in Electrical Engineering

Thesis: Directional Information Flow and Applications.

National Institute of Technology Karnataka, Surathkal *May 2008*

B.Tech. in Electrical and Electronics Engineering